

Logistic_Regression_Project (1)

December 9, 2024

Project Title

Stock price prediction system

Project Description

The aim of this project is to develop a stock price prediction system for Tesla that generates actionable buy or no-buy signals based on historical stock data. Specifically, the system will predict whether purchasing Tesla stock will be beneficial by analyzing past performance and trends.

Data Source: • The project uses historical stock data from Tesla, which includes key metrics such as the stock's opening, closing, high, and low prices, along with the trading volume.

Goal: • To predict a binary signal indicating whether the stock's closing price on the next trading day will be higher (buy signal) or lower (no-buy signal) compared to the current day's closing price.

Use Case: • The prediction system is designed to assist investors and traders by providing an automated recommendation on whether to buy Tesla stock. This recommendation is based on historical data patterns and trends.

Data Collection and Preparation: • Load the Tesla stock dataset and preprocess the data. This includes handling missing values, removing redundant columns, and preparing the data for analysis.

Feature Engineering: • Create new features that may enhance model performance. This includes calculating price differences (e.g., open-close) and adding time-based indicators such as quarter-end markers.

Exploratory Data Analysis (EDA): • Perform EDA to understand the dataset's structure, identify trends, and analyze correlations between features. This helps in gaining insights and preparing the data for modeling.

Model Building: • Develop machine learning models to classify the buy or no-buy signals. Algorithms such as Logistic Regression and Support Vector Machine (SVM) will be used to make these predictions.

Model Evaluation: • Assess the performance of the models using appropriate metrics (e.g., ROC-AUC) to determine which model provides the most reliable and accurate predictions. The best performing model will be selected based on its ability to generalize to unseen data.

Target Variable (For Supervised Learning)

1. Definition: The target variable is a binary indicator that represents whether it is potentially profitable to buy Tesla stock based on its predicted price movement from one day to the next.
2. How It's Calculated:

- 1 (Buy Signal): The target variable is set to 1 if the closing price of the Tesla stock on the next trading day is expected to be higher than the closing price on the current day. This prediction suggests a potential opportunity to buy the stock because the model forecasts a price increase.
- 0 (Don't Buy Signal): The target variable is set to 0 if the closing price of the Tesla stock on the next trading day is expected to be lower or the same as the closing price on the current day. This prediction suggests that it may not be advantageous to buy the stock because the model does not forecast a price increase.

3. How to Determine the Target Variable:

• To determine whether the target is 1 or 0, the model compares the closing price of today with the predicted closing price of the next trading day. • If `Next_Day_Close > Today_Close`, then the target is set to 1 (indicating a buy signal). • If `Next_Day_Close <= Today_Close`, then the target is set to 0 (indicating a no-buy signal). 4.

Use in Decision-Making: • The target variable helps in making stock trading decisions by predicting the direction of the stock price. A '1' (Buy Signal) indicates a potential opportunity to profit by purchasing the stock, while a '0' (Don't Buy Signal) suggests avoiding a purchase due to expected price stagnation or decline.

```
[169]: from google.colab import drive
drive.mount('/content/drive')
```

Drive already mounted at `/content/drive`; to attempt to forcibly remount, call `drive.mount("/content/drive", force_remount=True)`.

Import Libraries

```
[170]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

Machine Learning tools

```
[171]: from sklearn.model_selection import train_test_split, RandomizedSearchCV, \
↳ TimeSeriesSplit #Hyperparameter tuning (RandomSearchCV)
from sklearn.linear_model import LogisticRegression
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import (
    classification_report,
    precision_recall_curve,
    roc_curve,
    auc,
    confusion_matrix,
    roc_auc_score,
)
from imblearn.over_sampling import SMOTE
from imblearn.under_sampling import TomekLinks
```

```
from scipy.stats import uniform

import warnings
warnings.filterwarnings('ignore')
```

Load the dataset

```
[172]: df = pd.read_csv('/content/drive/My Drive/ML Assignment/TSLA.csv')
```

```
[173]: df.head()
```

```
[173]:
```

	Date	High	Low	Open	Close	Volume \
0	2019-09-30	48.796001	47.222000	48.599998	48.174000	29399000.0
1	2019-10-01	49.189999	47.826000	48.299999	48.938000	30813000.0
2	2019-10-02	48.930000	47.886002	48.658001	48.625999	28157000.0
3	2019-10-03	46.896000	44.855999	46.372002	46.605999	75422500.0
4	2019-10-04	46.956001	45.613998	46.321999	46.285999	39975000.0

	Adj Close
0	48.174000
1	48.938000
2	48.625999
3	46.605999
4	46.285999

```
[174]: df.tail()
```

```
[174]:
```

	Date	High	Low	Open	Close \
634	2022-04-05	1152.869995	1087.300049	1136.300049	1091.260010
635	2022-04-06	1079.000000	1027.699951	1073.469971	1045.760010
636	2022-04-07	1076.589966	1021.539978	1052.390015	1057.260010
637	2022-04-08	1048.439941	1022.440002	1043.209961	1025.489990
638	2022-04-11	1008.469971	974.640015	980.400024	975.929993

	Volume	Adj Close
634	26691700.0	1091.260010
635	29782800.0	1045.760010
636	26482400.0	1057.260010
637	18293300.0	1025.489990
638	19660500.0	975.929993

```
[175]: df.shape
```

```
[175]: (639, 7)
```

```
[176]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

RangeIndex: 639 entries, 0 to 638

Data columns (total 7 columns):

#	Column	Non-Null Count	Dtype
0	Date	639 non-null	object
1	High	639 non-null	float64
2	Low	639 non-null	float64
3	Open	639 non-null	float64
4	Close	639 non-null	float64
5	Volume	639 non-null	float64
6	Adj Close	639 non-null	float64

dtypes: float64(6), object(1)

memory usage: 35.1+ KB

```
[177]: df.describe()
```

```
[177]:
```

	High	Low	Open	Close	Volume \
count	639.000000	639.000000	639.000000	639.000000	6.390000e+02
mean	543.362885	517.883537	531.004088	531.298030	4.819130e+07
std	340.837426	325.395864	333.534448	333.362040	3.579030e+07
min	46.896000	44.855999	45.959999	46.285999	9.800600e+06
25%	170.258003	162.379997	167.349998	164.783005	2.392195e+07
50%	620.409973	595.500000	603.880005	605.130005	3.448900e+07
75%	796.584991	767.744995	779.445007	781.304993	6.329725e+07
max	1243.489990	1217.000000	1234.410034	1229.910034	3.046940e+08

	Adj Close
count	639.000000
mean	531.298030
std	333.362040
min	46.285999
25%	164.783005
50%	605.130005
75%	781.304993
max	1229.910034

```
[178]: df.columns = df.columns.str.strip() #Remove any leading or trailing white spaces
```

```
[179]: print("Columns:", df.columns)
```

Columns: Index(['Date', 'High', 'Low', 'Open', 'Close', 'Volume', 'Adj Close'], dtype='object')

Exploratory Data Analysis

```
[180]: plt.figure(figsize=(15,5))
plt.plot(df['Close'])
```

```
plt.title('Tesla Close price.', fontsize=15)
plt.ylabel('Price in dollars.')
plt.show()
```



```
[181]: df[df['Close'] == df['Adj Close']].shape
```

```
[181]: (639, 7)
```

```
[182]: df = df.drop(['Adj Close'], axis=1)  #Drop columns
```

Check missing values in each column

```
[183]: df.isnull().sum()
```

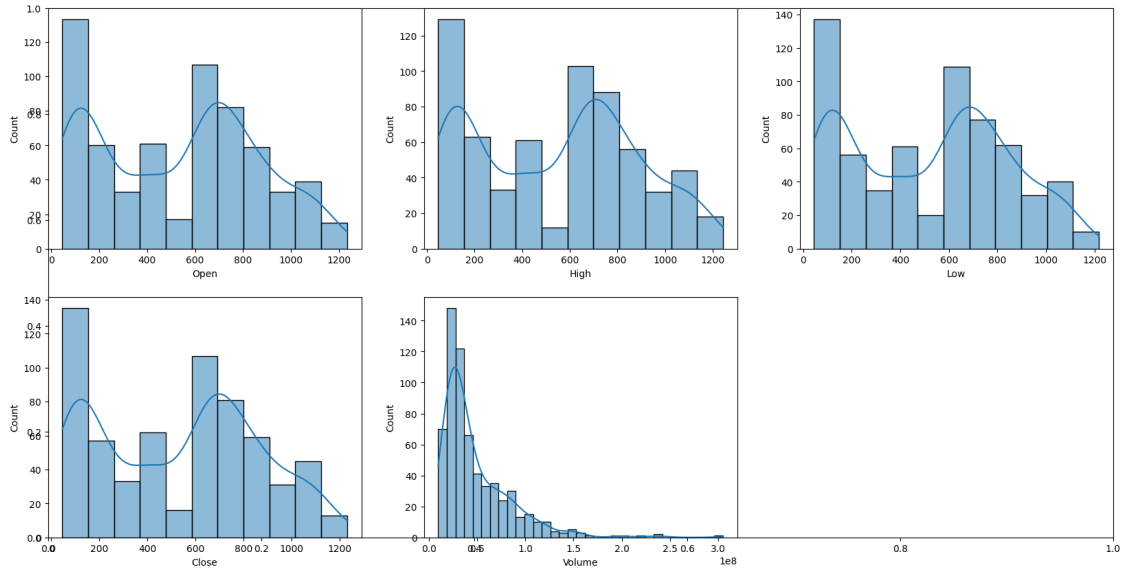
```
[183]: Date      0
      High      0
      Low       0
      Open      0
      Close     0
      Volume    0
      dtype: int64
```

Plot the Histogram

```
[184]: features = ['Open', 'High', 'Low', 'Close', 'Volume']

plt.subplots(figsize=(20,10))

for i, col in enumerate(features):
    plt.subplot(2,3,i+1)
    sns.histplot(df[col] , kde=True)
plt.show()
```



```
[185]: splitted = df['Date'].str.split('-', expand=True)

df['Year'] = splitted[0].astype('int')
df['Month'] = splitted[1].astype('int')
df['Day'] = splitted[2].astype('int')

df.head()
```

```
[185]:
```

	Date	High	Low	Open	Close	Volume	Year	\
0	2019-09-30	48.796001	47.222000	48.599998	48.174000	29399000.0	2019	
1	2019-10-01	49.189999	47.826000	48.299999	48.938000	30813000.0	2019	
2	2019-10-02	48.930000	47.886002	48.658001	48.625999	28157000.0	2019	
3	2019-10-03	46.896000	44.855999	46.372002	46.605999	75422500.0	2019	
4	2019-10-04	46.956001	45.613998	46.321999	46.285999	39975000.0	2019	

	Month	Day
0	9	30
1	10	1
2	10	2
3	10	3
4	10	4

```
[186]: df['is_quarter_end'] = np.where(df['Month']%3==0,1,0)
df.head()
```

```
[186]:
```

	Date	High	Low	Open	Close	Volume	Year	\
0	2019-09-30	48.796001	47.222000	48.599998	48.174000	29399000.0	2019	
1	2019-10-01	49.189999	47.826000	48.299999	48.938000	30813000.0	2019	


2	2019-10-02	48.930000	47.886002	48.658001	48.625999	28157000.0	2019
3	2019-10-03	46.896000	44.855999	46.372002	46.605999	75422500.0	2019
4	2019-10-04	46.956001	45.613998	46.321999	46.285999	39975000.0	2019

	Month	Day	is_quarter_end
0	9	30	1
1	10	1	0
2	10	2	0
3	10	3	0
4	10	4	0

Plot the mean value of the each column

```
[187]: data_grouped = df.drop('Date', axis=1).groupby('Year').mean()
plt.subplots(figsize=(20,10))

for i, col in enumerate(['Open', 'High', 'Low', 'Close']):
    plt.subplot(2,2,i+1)
    data_grouped[col].plot.bar()
plt.show()
```



```
[188]: df.drop('Date',axis=1).groupby('is_quarter_end').mean()
```


```
[188]:
```

	High	Low	Open	Close	Volume \
is_quarter_end					
0	537.826601	512.947180	525.494497	525.656104	4.886455e+07
1	553.906989	527.285055	541.497355	542.043335	4.690906e+07

	Year	Month	Day
is_quarter_end			
0	2020.508353	6.076372	15.527446
1	2020.509091	7.431818	15.750000

```
[189]: df['open-close'] = df['Open'] - df['Close']
df['low-high'] = df['Low'] - df['High']
df['target'] = np.where(df['Close'].shift(-1) > df['Close'], 1, 0) #Check next
    ↪ day closing price is greater than the current day closing price
```

```
[190]: plt.pie(df['target'].value_counts().values,
            labels=[0, 1], autopct='%1.1f%%')
plt.show()
```



```
[191]: df_reset = df.reset_index()
print(df_reset.columns)
```

```
Index(['index', 'Date', 'High', 'Low', 'Open', 'Close', 'Volume', 'Year',
      'Month', 'Day', 'is_quarter_end', 'open-close', 'low-high', 'target'],
      dtype='object')
```

```
[192]: df_reset = df.reset_index()
print(df_reset.head()) # Shows the first few rows of the DataFrame
```


	index	Date	High	Low	Open	Close	Volume	\
0	0	2019-09-30	48.796001	47.222000	48.599998	48.174000	29399000.0	
1	1	2019-10-01	49.189999	47.826000	48.299999	48.938000	30813000.0	
2	2	2019-10-02	48.930000	47.886002	48.658001	48.625999	28157000.0	
3	3	2019-10-03	46.896000	44.855999	46.372002	46.605999	75422500.0	
4	4	2019-10-04	46.956001	45.613998	46.321999	46.285999	39975000.0	

	Year	Month	Day	is_quarter_end	open-close	low-high	target
0	2019	9	30	1	0.425999	-1.574001	1
1	2019	10	1	0	-0.638000	-1.363998	0
2	2019	10	2	0	0.032001	-1.043999	0
3	2019	10	3	0	-0.233997	-2.040001	0
4	2019	10	4	0	0.035999	-1.342003	1

```
[193]: print(df_reset.index)
```

```
RangeIndex(start=0, stop=639, step=1)
```

```
[194]: print(df.head()) # Check the first few rows
print(df.columns) # Ensure 'Date' is listed
```

	Date	High	Low	Open	Close	Volume	Year	\
0	2019-09-30	48.796001	47.222000	48.599998	48.174000	29399000.0	2019	
1	2019-10-01	49.189999	47.826000	48.299999	48.938000	30813000.0	2019	
2	2019-10-02	48.930000	47.886002	48.658001	48.625999	28157000.0	2019	
3	2019-10-03	46.896000	44.855999	46.372002	46.605999	75422500.0	2019	
4	2019-10-04	46.956001	45.613998	46.321999	46.285999	39975000.0	2019	

	Month	Day	is_quarter_end	open-close	low-high	target
0	9	30	1	0.425999	-1.574001	1
1	10	1	0	-0.638000	-1.363998	0
2	10	2	0	0.032001	-1.043999	0
3	10	3	0	-0.233997	-2.040001	0
4	10	4	0	0.035999	-1.342003	1

```
Index(['Date', 'High', 'Low', 'Open', 'Close', 'Volume', 'Year', 'Month',
      'Day', 'is_quarter_end', 'open-close', 'low-high', 'target'],
      dtype='object')
```

```
[195]: print([col for col in df.columns if 'Date' in col.lower()]) # Checks for
↳ similar names
```


```
[]
```

```
[196]: # Additional Heatmap for Highly Correlated Features
plt.figure(figsize=(10, 10))

# Reset the index temporarily to drop 'Date' column
df_reset = df.reset_index()
```

```
# Generate the heatmap with correlation greater than 0.9
sns.heatmap(df_reset.drop('Date', axis=1).corr() > 0.9, annot=True, cbar=False)

plt.title('Heatmap for Features with Correlation > 0.9')
plt.show()
```



```
[197]: features = df[['open-close', 'low-high', 'is_quarter_end']]
target = df['target']

scaler = StandardScaler()
features = scaler.fit_transform(features)
```

```
X_train, X_valid, Y_train, Y_valid = train_test_split(
    features, target, test_size=0.1, random_state=2022)
print(X_train.shape, X_valid.shape)
```

(575, 3) (64, 3)

Feature Engineering

```
[198]: df['Price Change'] = df['Close'] - df['Open']
df['Price Direction'] = np.where(df['Price Change'] > 0, 1, 0)
```

Adding lagged features for temporal patterns

```
[199]: for lag in range(1, 4): # 1-day, 2-day, 3-day lag
    df[f'Lag_Close_{lag}'] = df['Close'].shift(lag)    #Learn from past details
```

Drop rows with Nan values after adding lagged features

```
[200]: df.dropna(inplace=True)
```

Select features and target

```
[201]: X = df[['High', 'Low', 'Open', 'Close', 'Lag_Close_1', 'Lag_Close_2', 'Lag_Close_3']]
y = df['Price Direction']
```

Scale features

```
[202]: scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)
```

Train-Test Split

```
[203]: train_size = int(len(df) * 0.7)
valid_size = int(len(df) * 0.85)
```

```
[204]: X_train, X_valid, X_test = (
    X_scaled[:train_size],
    X_scaled[train_size:valid_size],
    X_scaled[valid_size:],
)
```

```
[205]: y_train, y_valid, y_test = (
    y[:train_size].values,
    y[train_size:valid_size].values,
    y[valid_size:].values,
)
```

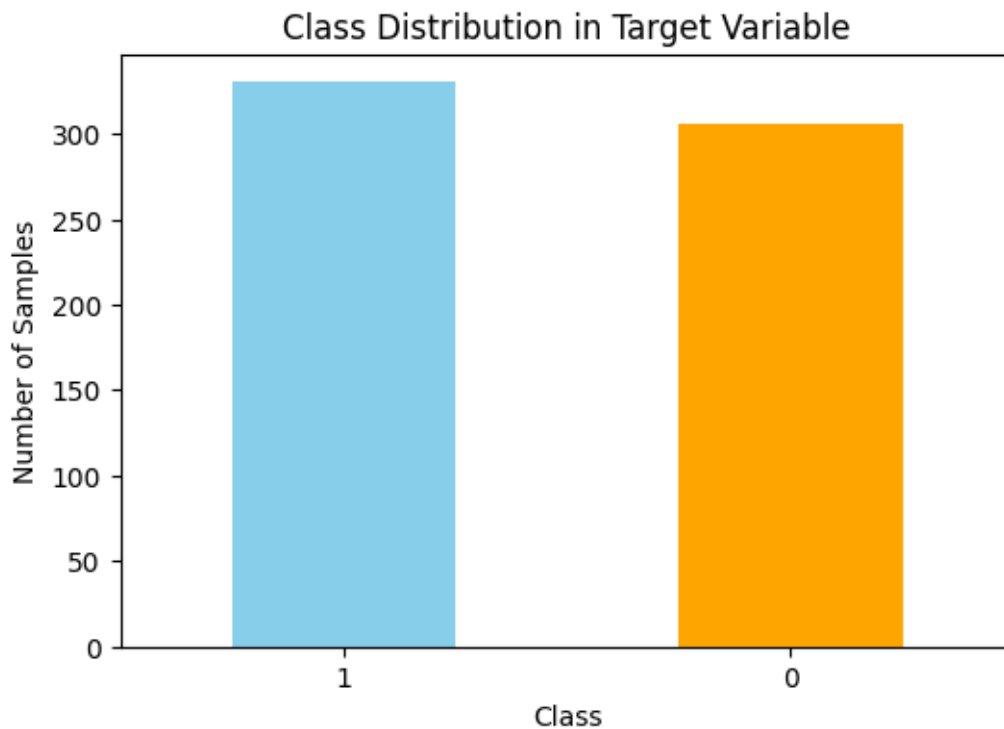
Check the class distribution

```
[206]: class_counts = df['Price Direction'].value_counts()
print("Class Distribution:")
print(class_counts)
```

```
Class Distribution:
Price Direction
1      330
0      306
Name: count, dtype: int64
```

Visualize the class distribution

```
[207]: plt.figure(figsize=(6, 4))
class_counts.plot(kind='bar', color=['skyblue', 'orange'])
plt.title("Class Distribution in Target Variable")
plt.xlabel("Class")
plt.ylabel("Number of Samples")
plt.xticks(rotation=0)
plt.show()
```



Calculate the imbalanced ratio

```
[208]: imbalance_ratio = class_counts.max() / class_counts.min()
print(f"\nImbalance Ratio: {imbalance_ratio:.2f}")
```

Imbalance Ratio: 1.08

```
[209]: if imbalance_ratio > 1.5:
        print("The dataset is imbalanced.")
    else:
        print("The dataset is balanced.")
```

The dataset is balanced.

Handle imbalance using SMOTE and Tomek Links

```
[210]: smote = SMOTE(random_state=42)
        X_train_res, y_train_res = smote.fit_resample(X_train, y_train)
```

```
[211]: tomek = TomekLinks()
        X_train_res, y_train_res = tomek.fit_resample(X_train_res, y_train_res)
        ↪ #Further cleaning resampling dataset
```

Logistic regression with Hyperparameter tuning

```
[212]: param_dist = {
        'C': uniform(0.1, 10),    #Regularization parameter
        'penalty': ['l2'],        #Regularization type
        'solver': ['lbfgs'],      #Algorithm for optimization
        'max_iter': [100, 200, 500],
        'class_weight': ['balanced', None],
    }
```

TimeSeries Split for stability

```
[213]: tscv = TimeSeriesSplit(n_splits=5)    #Ensure trainig data always comes before
        ↪ the validation data

        random_search = RandomizedSearchCV(
            LogisticRegression(),
            param_distributions=param_dist,
            n_iter=50,
            cv=tscv,
            scoring='roc_auc',
            random_state=42,
            n_jobs=-1,
        )
        random_search.fit(X_train_res, y_train_res)
```

```
[213]: RandomizedSearchCV(cv=TimeSeriesSplit(gap=0, max_train_size=None, n_splits=5,
        test_size=None),
                               estimator=LogisticRegression(), n_iter=50, n_jobs=-1,
                               param_distributions={'C':
```

```
<scipy.stats._distn_infrastructure.rv_continuous_frozen object at
0x78690c29e290>,
                                'class_weight': ['balanced', None],
                                'max_iter': [100, 200, 500],
                                'penalty': ['l2'],
                                'solver': ['lbfgs']},
                                random_state=42, scoring='roc_auc')
```

Best model

```
[214]: best_model = random_search.best_estimator_
print("\nBest Parameters:", random_search.best_params_)
```

```
Best Parameters: {'C': 9.75255307264138, 'class_weight': None, 'max_iter': 100,
'penalty': 'l2', 'solver': 'lbfgs'}
```

Train best model

```
[215]: best_model.fit(X_train_res, y_train_res)
```

```
[215]: LogisticRegression(C=9.75255307264138)
```

Validation Predictions

```
[216]: y_pred_proba_valid = best_model.predict_proba(X_valid)[:, 1]
```

Optimize threshold based on Precision-Recall curve

```
[217]: precision, recall, thresholds = precision_recall_curve(y_valid,
    ↪ y_pred_proba_valid)
optimal_idx = np.argmax(precision * recall)
optimal_threshold = thresholds[optimal_idx]
print(f"\nOptimal Threshold: {optimal_threshold:.2f}")

y_pred_valid = (y_pred_proba_valid >= optimal_threshold).astype(int)
```

Optimal Threshold: 0.43

Classification Report (Validation)

```
[218]: print("\nValidation Classification Report:")
print(classification_report(y_valid, y_pred_valid))
```

Validation Classification Report:

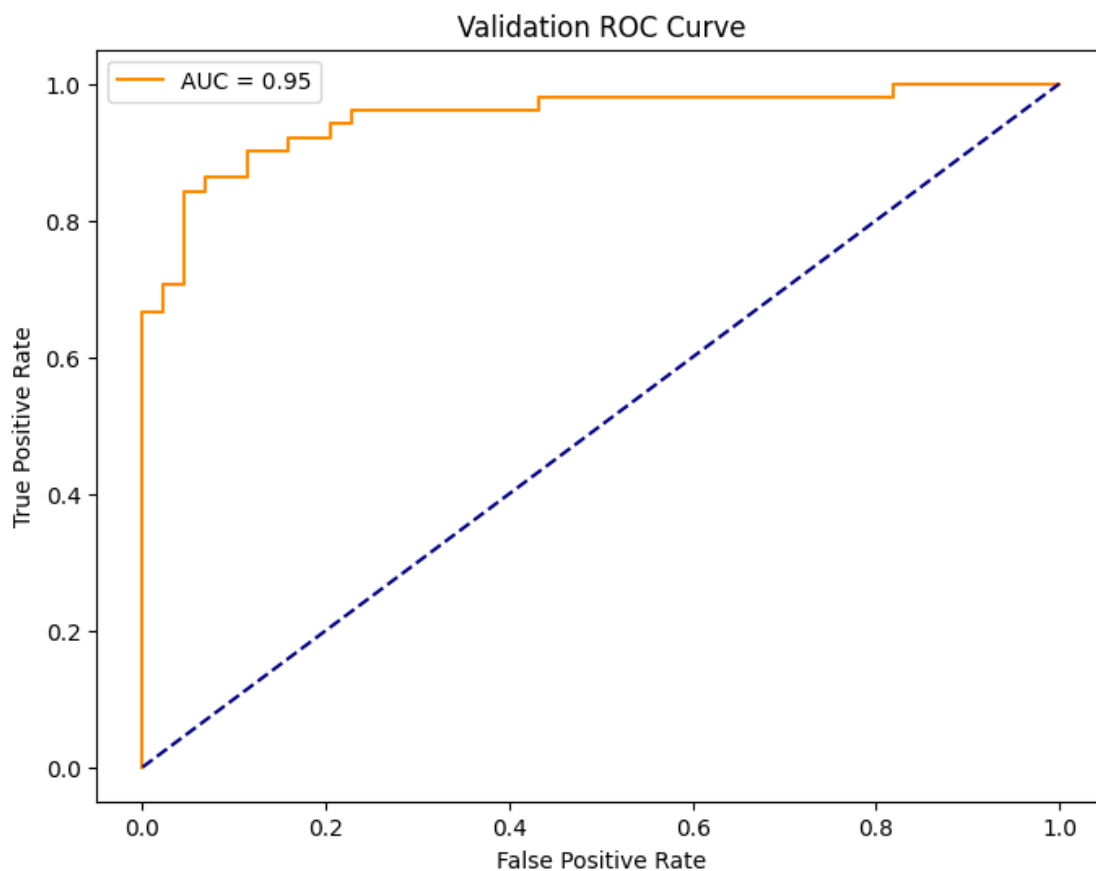
	precision	recall	f1-score	support
0	0.89	0.89	0.89	44

	1	0.90	0.90	0.90	51
accuracy				0.89	95
macro avg		0.89	0.89	0.89	95
weighted avg		0.89	0.89	0.89	95

ROC Curve (Validation)

```
[219]: fpr, tpr, _ = roc_curve(y_valid, y_pred_proba_valid)
roc_auc = auc(fpr, tpr)

plt.figure(figsize=(8, 6))
plt.plot(fpr, tpr, label=f'AUC = {roc_auc:.2f}', color='darkorange')
plt.plot([0, 1], [0, 1], color='navy', linestyle='--')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Validation ROC Curve')
plt.legend(loc='best')
plt.show()
```



Final Test Evaluation

```
[220]: y_pred_proba_test = best_model.predict_proba(X_test)[: , 1]
y_pred_test = (y_pred_proba_test >= optimal_threshold).astype(int)

test_auc = roc_auc_score(y_test, y_pred_proba_test)
print(f"\nTest AUC: {test_auc:.4f}")

print("\nTest Classification Report:")
print(classification_report(y_test, y_pred_test))
```

Test AUC: 0.9961

Test Classification Report:

	precision	recall	f1-score	support
0	0.88	1.00	0.94	45
1	1.00	0.88	0.94	51
accuracy			0.94	96
macro avg	0.94	0.94	0.94	96
weighted avg	0.94	0.94	0.94	96

Confusion matrix (Test set)

```
[221]: y_pred_proba_test = best_model.predict_proba(X_test)[: , 1]
y_pred_test = (y_pred_proba_test >= optimal_threshold).astype(int)

test_auc = roc_auc_score(y_test, y_pred_proba_test)
print(f"\nTest AUC: {test_auc:.4f}")

print("\nTest Classification Report:")
print(classification_report(y_test, y_pred_test))
```

Test AUC: 0.9961


Test Classification Report:

	precision	recall	f1-score	support
0	0.88	1.00	0.94	45
1	1.00	0.88	0.94	51
accuracy			0.94	96
macro avg	0.94	0.94	0.94	96
weighted avg	0.94	0.94	0.94	96

ROC Curve (Test set)

```
[222]: fpr_test, tpr_test, _ = roc_curve(y_test, y_pred_proba_test)
roc_auc_test = auc(fpr_test, tpr_test)

plt.figure(figsize=(8, 6))
plt.plot(fpr_test, tpr_test, label=f'AUC = {roc_auc_test:.2f}',
        color='darkorange')
plt.plot([0, 1], [0, 1], color='navy', linestyle='--')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Test ROC Curve')
plt.legend(loc='best')
plt.show()
```



Precision-Recall Curve (Test Set)

```
[223]: precision_test, recall_test, _ = precision_recall_curve(y_test,
        y_pred_proba_test)
plt.figure(figsize=(8, 6))
```

```
plt.plot(recall_test, precision_test, label=f'AP = {auc(recall_test, precision_test):.2f}')
plt.xlabel('Recall')
plt.ylabel('Precision')
plt.title('Test Precision-Recall Curve')
plt.legend(loc='best')
plt.show()
```

