

SVM_Project (1)

December 9, 2024

```
[196]: from google.colab import drive
drive.mount('/content/drive')
```

Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive", force_remount=True).

Import Libraries

```
[197]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

Machine Learning tools

```
[198]: from sklearn.preprocessing import StandardScaler
from sklearn.svm import SVC
from sklearn.model_selection import train_test_split, RandomizedSearchCV, \
    TimeSeriesSplit, cross_val_score
from sklearn.metrics import classification_report, roc_curve, auc, \
    precision_recall_curve, confusion_matrix
from imblearn.over_sampling import SMOTE
from imblearn.under_sampling import TomekLinks
from sklearn import metrics
from scipy.stats import uniform, randint
from sklearn.decomposition import PCA

import warnings
warnings.filterwarnings('ignore')
```

```
[199]: df = pd.read_csv('/content/drive/My Drive/ML Assignment/TSLA.csv')
df.head()
```

```
[199]:
```

	Date	High	Low	Open	Close	Volume	\
0	2019-09-30	48.796001	47.222000	48.599998	48.174000	29399000.0	
1	2019-10-01	49.189999	47.826000	48.299999	48.938000	30813000.0	
2	2019-10-02	48.930000	47.886002	48.658001	48.625999	28157000.0	
3	2019-10-03	46.896000	44.855999	46.372002	46.605999	75422500.0	

```
4 2019-10-04 46.956001 45.613998 46.321999 46.285999 39975000.0
```

```
Adj Close
0 48.174000
1 48.938000
2 48.625999
3 46.605999
4 46.285999
```

```
[200]: df.tail()
```

```
[200]:
```

	Date	High	Low	Open	Close \
634	2022-04-05	1152.869995	1087.300049	1136.300049	1091.260010
635	2022-04-06	1079.000000	1027.699951	1073.469971	1045.760010
636	2022-04-07	1076.589966	1021.539978	1052.390015	1057.260010
637	2022-04-08	1048.439941	1022.440002	1043.209961	1025.489990
638	2022-04-11	1008.469971	974.640015	980.400024	975.929993

	Volume	Adj Close
634	26691700.0	1091.260010
635	29782800.0	1045.760010
636	26482400.0	1057.260010
637	18293300.0	1025.489990
638	19660500.0	975.929993

```
[201]: df.shape
```

```
[201]: (639, 7)
```

```
[202]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 639 entries, 0 to 638
Data columns (total 7 columns):
#   Column      Non-Null Count  Dtype
---  -
0   Date        639 non-null   object
1   High        639 non-null   float64
2   Low         639 non-null   float64
3   Open        639 non-null   float64
4   Close       639 non-null   float64
5   Volume      639 non-null   float64
6   Adj Close   639 non-null   float64
dtypes: float64(6), object(1)
memory usage: 35.1+ KB
```

```
[203]: df.describe()
```

```
[203]:
```

	High	Low	Open	Close	Volume \
count	639.000000	639.000000	639.000000	639.000000	6.390000e+02
mean	543.362885	517.883537	531.004088	531.298030	4.819130e+07
std	340.837426	325.395864	333.534448	333.362040	3.579030e+07
min	46.896000	44.855999	45.959999	46.285999	9.800600e+06
25%	170.258003	162.379997	167.349998	164.783005	2.392195e+07
50%	620.409973	595.500000	603.880005	605.130005	3.448900e+07
75%	796.584991	767.744995	779.445007	781.304993	6.329725e+07
max	1243.489990	1217.000000	1234.410034	1229.910034	3.046940e+08

	Adj Close
count	639.000000
mean	531.298030
std	333.362040
min	46.285999
25%	164.783005
50%	605.130005
75%	781.304993
max	1229.910034

Feature engineering

```
[204]: df['Price Change'] = df['Close'] - df['Open']
df['Price Direction'] = np.where(df['Price Change'] > 0, 1, 0)
```

```
[205]: X = df[['High', 'Low', 'Open', 'Close', 'Volume']]
y = df['Price Direction']
```

Scaling features

```
[206]: scaler = StandardScaler()
```

```
[207]: X_scaled = scaler.fit_transform(X)
```

Train-validation-test split

```
[208]: train_size = int(len(df) * 0.7) # 70% training
valid_size = int(len(df) * 0.15) # 15% validation
test_size = len(df) - train_size - valid_size
```

```
[209]: X_train, X_valid, X_test = X_scaled[:train_size], X_scaled[train_size:
    ↪train_size + valid_size], X_scaled[train_size + valid_size:]
y_train, y_valid, y_test = y[:train_size], y[train_size:train_size +
    ↪valid_size], y[train_size + valid_size:]
```


Check the class distribution

```
[210]: class_counts = df['Price Direction'].value_counts()
print("Class Distribution:")
print(class_counts)
```

```
Class Distribution:
Price Direction
1      331
0      308
Name: count, dtype: int64
```

Visualize the class distribution

```
[211]: plt.figure(figsize=(6, 4))
class_counts.plot(kind='bar', color=['skyblue', 'orange'])
plt.title("Class Distribution in Target Variable")
plt.xlabel("Class")
plt.ylabel("Number of Samples")
plt.xticks(rotation=0)
plt.show()
```



Calculate the imbalanced ratio

```
[212]: imbalance_ratio = class_counts.max() / class_counts.min()
print(f"\nImbalance Ratio: {imbalance_ratio:.2f}")
```

Imbalance Ratio: 1.07

```
[213]: if imbalance_ratio > 1.5:
        print("The dataset is imbalanced.")
    else:
        print("The dataset is balanced.")
```

The dataset is balanced.

Apply SMOTE to balance the classes in the training set

```
[214]: smote = SMOTE(random_state=42)
        X_train_res, y_train_res = smote.fit_resample(X_train, y_train)
```

Apply Tomek Links to remove noise and refine boundaries

```
[215]: tomek = TomekLinks()
        X_train_res, y_train_res = tomek.fit_resample(X_train_res, y_train_res)
```

Hyperparameter tuning using RandomizedSearchCV for SVM

```
[216]: param_dist = {
        'C': uniform(0.1, 10), # generate random values between 0.1 and 10
        'gamma': ['scale', 'auto', 0.1, 0.5, 1], # Gamma values for kernel
        'kernel': ['linear'], # Different kernel types
        'class_weight': ['balanced', None],
        'degree': [3], # Degree for polynomial kernel
        'tol': [1e-3, 1e-4] # Tolerance for stopping criterion
    }
```

Time series split for time series validation

```
[217]: tscv = TimeSeriesSplit(n_splits=3) #Avoid data leakage
```

Perform RandomizedSearchCV for SVM

```
[218]: random_search = RandomizedSearchCV(SVC(probability=True),
        ↪param_distributions=param_dist, n_iter=100, cv=tscv, scoring='roc_auc',
        ↪random_state=42, n_jobs=-1)
        random_search.fit(X_train_res, y_train_res)
```

```
[218]: RandomizedSearchCV(cv=TimeSeriesSplit(gap=0, max_train_size=None, n_splits=3,
        test_size=None),
        estimator=SVC(probability=True), n_iter=100, n_jobs=-1,
        param_distributions={'C':
        <scipy.stats._distn_infrastructure.rv_continuous_frozen object at
        0x7bae2275f850>,
        'class_weight': ['balanced', None],
```

```

        'degree': [3],
        'gamma': ['scale', 'auto', 0.1, 0.5, 1],
        'kernel': ['linear'],
        'tol': [0.001, 0.0001]},
    random_state=42, scoring='roc_auc')

```

Best model from RandomizedSearchCV

```

[219]: best_model_svm = random_search.best_estimator_
print("\nBest Parameters from RandomizedSearchCV:", random_search.best_params_)

```

```

Best Parameters from RandomizedSearchCV: {'C': 9.494989415641891,
'class_weight': None, 'degree': 3, 'gamma': 0.5, 'kernel': 'linear', 'tol':
0.0001}

```

Train the best SVM model

```

[220]: best_model_svm.fit(X_train_res, y_train_res)

```

```

[220]: SVC(C=9.494989415641891, gamma=0.5, kernel='linear', probability=True,
tol=0.0001)

```

Evaluate model on validation set

```

[221]: y_pred_svm = best_model_svm.predict(X_valid)

```

```

[222]: y_pred_proba_svm = best_model_svm.predict_proba(X_valid)[:, 1]

```

Evaluate model on test set

```

[223]: y_test_pred = best_model_svm.predict(X_test)
y_test_proba = best_model_svm.predict_proba(X_test)[:, 1]

```

Print classification report

```

[224]: print("\nClassification Report (Validation Data):")
print(classification_report(y_valid, y_pred_svm))

```

Classification Report (Validation Data):

	precision	recall	f1-score	support
0	0.67	0.95	0.79	43
1	0.94	0.62	0.74	52
accuracy			0.77	95
macro avg	0.81	0.78	0.77	95
weighted avg	0.82	0.77	0.76	95

```
[225]: print("\nClassification Report (Test Data):")
print(classification_report(y_test, y_test_pred))
```

```
Classification Report (Test Data):
```

	precision	recall	f1-score	support
0	0.98	0.98	0.98	46
1	0.98	0.98	0.98	51
accuracy			0.98	97
macro avg	0.98	0.98	0.98	97
weighted avg	0.98	0.98	0.98	97

Reduce to 3 Principal components

```
[226]: pca = PCA(n_components=3)    #For dimensionality reduction and faster training
```


```
[227]: X_scaled_pca = pca.fit_transform(X_scaled)
```

Compute AUC and plot ROC curve

```
[228]: fpr_valid, tpr_valid, _ = roc_curve(y_valid, y_pred_proba_svm)
fpr_test, tpr_test, _ = roc_curve(y_test, y_test_proba)

roc_auc_valid = auc(fpr_valid, tpr_valid)
roc_auc_test = auc(fpr_test, tpr_test)


plt.figure(figsize=(8, 6))
plt.plot(fpr_valid, tpr_valid, label=f'Validation AUC = {roc_auc_valid:.2f}',
        color='blue')
plt.plot(fpr_test, tpr_test, label=f'Test AUC = {roc_auc_test:.2f}',
        color='darkorange')
plt.plot([0, 1], [0, 1], color='navy', linestyle='--')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curve')
plt.legend(loc='best')
plt.show()
```



Plot precision-Recall curve


```
[229]: precision, recall, _ = precision_recall_curve(y_valid, y_pred_proba_svm)

plt.figure(figsize=(8, 6))
plt.plot(recall, precision, label=f'AUC = {auc(recall, precision):.2f}',
         color='darkgreen')
plt.xlabel('Recall')
plt.ylabel('Precision')
plt.title('Precision-Recall Curve')
plt.legend(loc='best')
plt.show()
```


```
[230]: precision_test, recall_test, _ = precision_recall_curve(y_test, y_test_proba)

plt.figure(figsize=(8, 6))
plt.plot(recall_test, precision_test, label=f'Test AUC = {auc(recall_test, precision_test):.2f}', color='darkgreen')
plt.xlabel('Recall')
plt.ylabel('Precision')
plt.title('Precision-Recall Curve (Test Data)')
plt.legend(loc='best')
plt.show()
```




Confusion Matrix for validation data

```
[231]: conf_matrix = confusion_matrix(y_valid, y_pred_svm)
plt.figure(figsize=(6, 6))
sns.heatmap(conf_matrix, annot=True, fmt='d', cmap='Blues', xticklabels=[0, 1],
            yticklabels=[0, 1])
plt.xlabel('Predicted')
plt.ylabel('True')
plt.title('Confusion Matrix for validation data')
plt.show()
```



```
[232]: conf_matrix_test = confusion_matrix(y_test, y_test_pred)
plt.figure(figsize=(6, 6))
sns.heatmap(conf_matrix_test, annot=True, fmt='d', cmap='Blues',
            xticklabels=[0, 1], yticklabels=[0, 1])
plt.xlabel('Predicted')
plt.ylabel('True')
plt.title('Confusion Matrix (Test Data)')
plt.show()
```



Cross-validation with SVM

```
[233]: cv_scores_res_svm = cross_val_score(best_model_svm, X_train_res, y_train_res,
      ↳ cv=tscv, scoring='roc_auc')
      print(f"\nCross-Validation AUC After SMOTE and Tomek Links: {cv_scores_res_svm.
      ↳ mean():.4f} ± {cv_scores_res_svm.std():.4f}")
```

Cross-Validation AUC After SMOTE and Tomek Links: 0.8491 ± 0.2044

Evaluate model on original data

```
[234]: train_auc = metrics.roc_auc_score(y_train_res, best_model_svm.
      ↳ predict_proba(X_train_res)[: , 1])
      valid_auc = metrics.roc_auc_score(y_valid, y_pred_proba_svm)
      test_auc = metrics.roc_auc_score(y_test, y_test_proba)
```

```
print(f"\nSVM Model Performance (AUC):")
print(f"Training AUC: {train_auc:.4f}")
print(f"Validation AUC: {valid_auc:.4f}")
print(f"Test AUC: {test_auc:.4f}")
```

SVM Model Performance (AUC):

Training AUC: 0.9045

Validation AUC: 0.9602

Test AUC: 0.9991

Comparison

1.AUC:

SVM has a significantly higher Test AUC (0.9991) compared to Logistic Regression (0.9961). Higher AUC suggests that SVM is better at distinguishing between classes across all threshold values.

2.Accuracy:

*SVM achieves higher Test Accuracy (0.98) than Logistic Regression Test Accuracy (0.94).

3.Precision, Recall, F1-Score:

*SVM has consistently higher scores (0.98) compared to Logistic Regression (0.94) on Test Data.

4.Validation Performance:

*Logistic Regression performs better in Validation Classification Report (Accuracy: 0.89) compared to SVM (Accuracy: 0.77).

*Logistic Regression shows better balance between classes in Validation Data.

5.Overfitting:

*SVM shows signs of potential overfitting since its performance drops on Validation Data (AUC: 0.9602) compared to Test Data (AUC: 0.9991). Logistic Regression appears more stable between Validation and Test Data.

Conclusion

*Test Performance is Critical

SVM is preferable, as it has higher AUC, Accuracy, and better classification metrics on the Test Data.

*Validation Performance is Important

Logistic Regression is more balanced, especially for unseen data (Validation Set), and less prone to overfitting.

Summary

Based on the given matrices SVM is better model when we prioritize overall test performance and higher AUC.

But according to my model, SVM may be overfitting because of that I think Logistic regression is better.